

Visualization (Exploring Co-variation)

AUTHOR
Peter Ganong and Maggie Shi

PUBLISHED
October 14, 2024

DataTransformerRegistry.enable('default')

Table of contents

- 1. Categorical variable and a continuous variable
- 2. Two categorical variables
- 3. Two continuous variables
- 4. Graphics for production

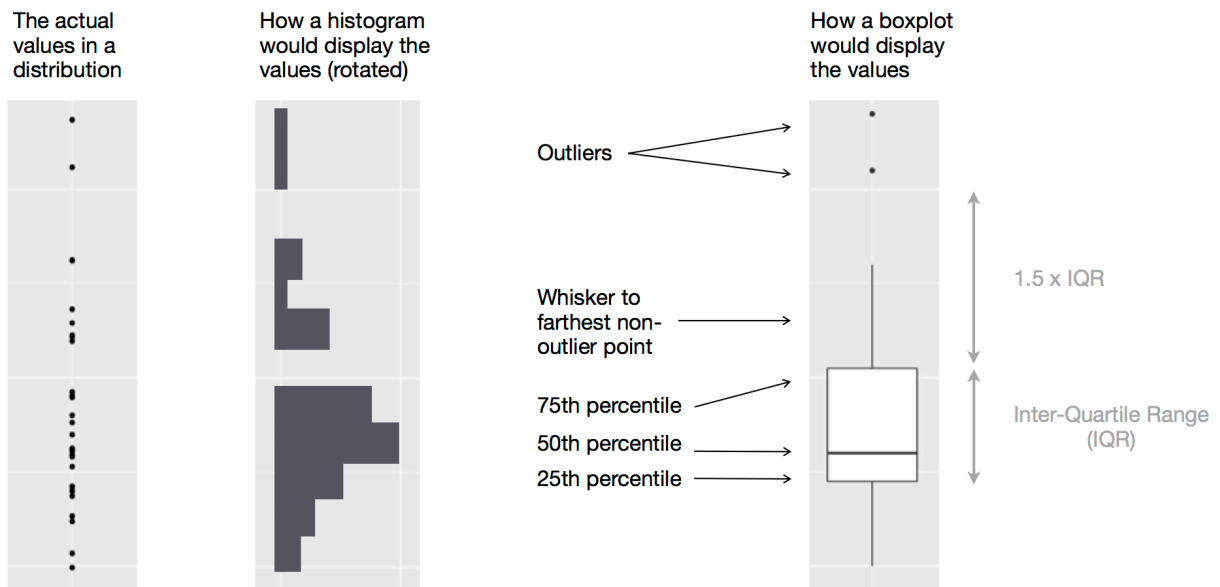
Categorical variable and continuous variable

```
from palmerpenguins import load_penguins
penguins = load_penguins()
display(penguins)
```

| | species | island | bill_length_mm | bill_depth_mm | flipper_length_mm | body_mass_g | sex | year |
|-----|-----------|-----------|----------------|---------------|-------------------|-------------|--------|------|
| 0 | Adelie | Torgersen | 39.1 | 18.7 | 181.0 | 3750.0 | male | 2007 |
| 1 | Adelie | Torgersen | 39.5 | 17.4 | 186.0 | 3800.0 | female | 2007 |
| 2 | Adelie | Torgersen | 40.3 | 18.0 | 195.0 | 3250.0 | female | 2007 |
| 3 | Adelie | Torgersen | NaN | NaN | NaN | NaN | NaN | 2007 |
| 4 | Adelie | Torgersen | 36.7 | 19.3 | 193.0 | 3450.0 | female | 2007 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 339 | Chinstrap | Dream | 55.8 | 19.8 | 207.0 | 4000.0 | male | 2009 |
| 340 | Chinstrap | Dream | 43.5 | 18.1 | 202.0 | 3400.0 | female | 2009 |
| 341 | Chinstrap | Dream | 49.6 | 18.2 | 193.0 | 3775.0 | male | 2009 |
| 342 | Chinstrap | Dream | 50.8 | 19.0 | 210.0 | 4100.0 | male | 2009 |
| 343 | Chinstrap | Dream | 50.2 | 18.7 | 198.0 | 3775.0 | female | 2009 |


344 rows × 8 columns

numeric & categorical: box plot



numeric & categorical: `mark_boxplot()`


```
alt.Chart(penguins).mark_boxplot().encode(
  x=alt.X('species:N', title="Species"),
  y=alt.Y('body_mass_g:Q', title="Body Mass (g)"),
).properties(
  width=400,
  height=300
)
```



Discussion question: what do you notice from this graph?

numeric & categorical: `transform_density()`

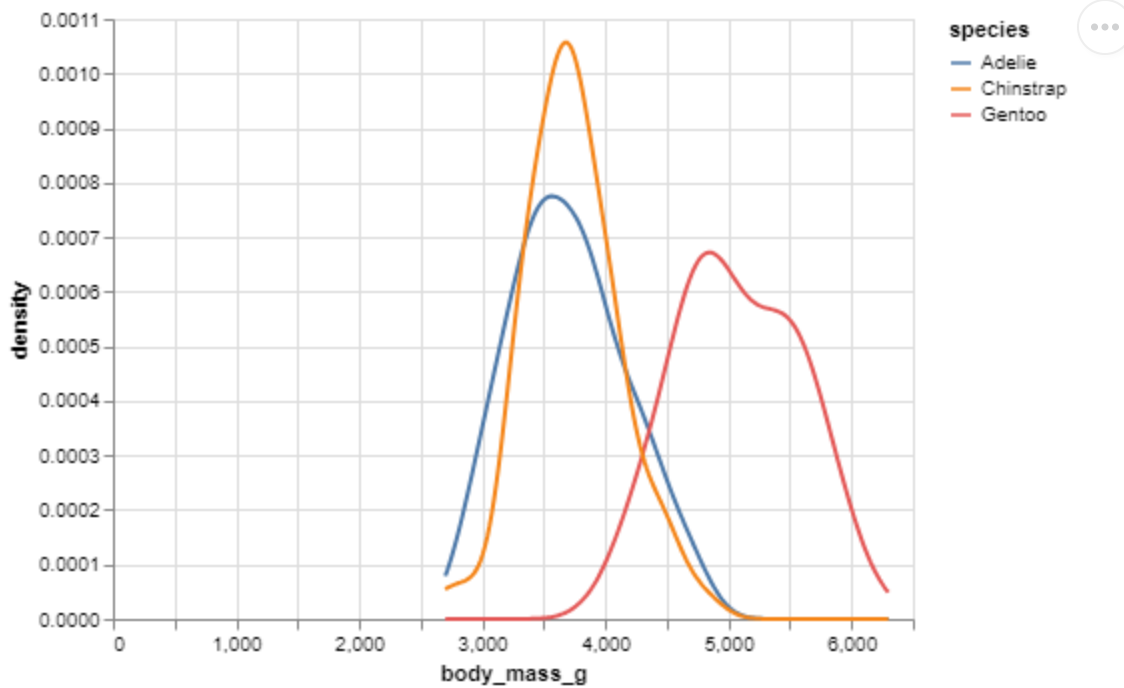
```
alt.Chart(penguins).transform_density(  
    'body_mass_g',  
    groupby=['species'],  
    as_=['body_mass_g', 'density']  
)  
.mark_line().encode(  
    alt.X('body_mass_g:Q'),  
    alt.Y('density:Q', stack=None),  
    alt.Color('species:N')  
)  
.properties(width=400,height=300)
```



numeric & categorical: `transform_density()`

Discussion q – What if we required the x-axis range to include zero? Would that improve or reduce clarity? How come?

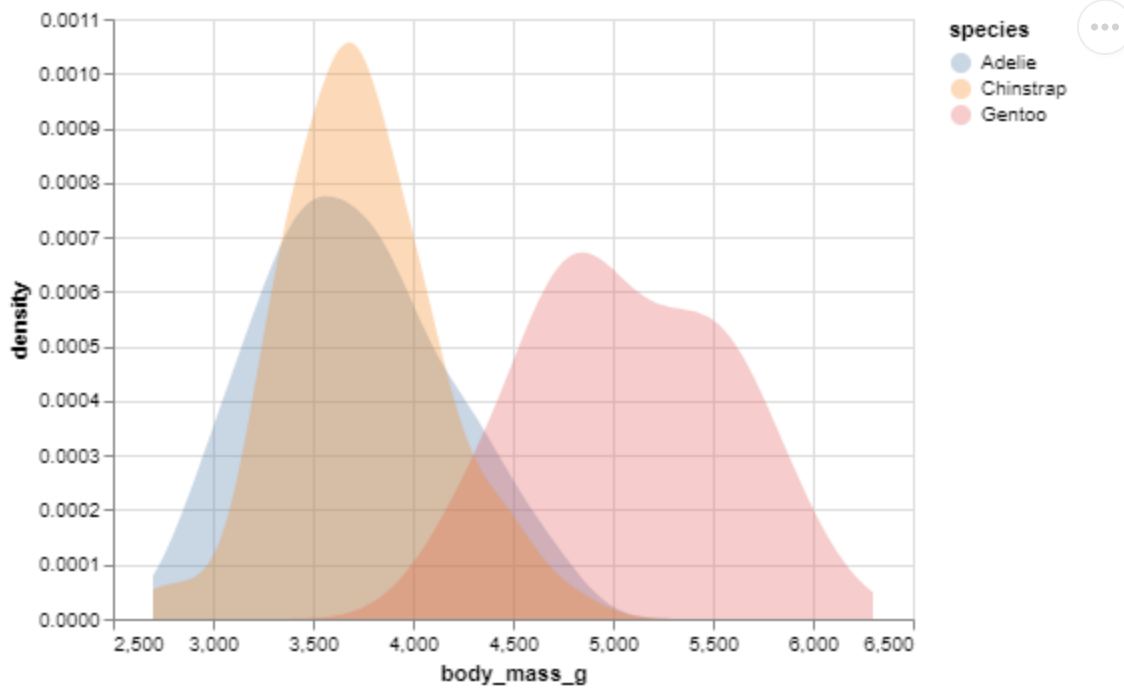
```
alt.Chart(penguins).transform_density(
    'body_mass_g',
    groupby=['species'],
    as_=['body_mass_g', 'density']
).mark_line().encode(
    alt.X('body_mass_g:Q', scale=alt.Scale(zero=True)),
    alt.Y('density:Q', stack=None),
    alt.Color('species:N')
).properties(width=400,height=300)
```



numeric & categorical: `transform_density()` filled in

`opacity=0.3` makes no difference in content; maybe a bit more elegant

```
alt.Chart(penguins).transform_density(
    'body_mass_g',
    groupby=['species'], # Group by species for different density curves
    as_=['body_mass_g', 'density']
).mark_area(opacity=0.3).encode(
    alt.X('body_mass_g:Q'),
    alt.Y('density:Q', stack=None),
    alt.Color('species:N')
).properties(width=400,height=300)
```



Two categorical variables

Question: How is cut related to color? 2 categorical vars


```
diamonds_grouped = diamonds.groupby(['color', 'cut']).size().reset_index().rename(columns={0: 'N'})
diamonds_grouped
```

| | color | cut | N |
|----|-------|-----------|------|
| 0 | D | Fair | 163 |
| 1 | D | Good | 662 |
| 2 | D | Very Good | 1513 |
| 3 | D | Premium | 1603 |
| 4 | D | Ideal | 2834 |
| 5 | E | Fair | 224 |
| 6 | E | Good | 933 |
| 7 | E | Very Good | 2400 |
| 8 | E | Premium | 2337 |
| 9 | E | Ideal | 3903 |
| 10 | F | Fair | 312 |
| 11 | F | Good | 909 |
| 12 | F | Very Good | 2164 |
| 13 | F | Premium | 2331 |

| | color | cut | N |
|----|-------|-----------|------|
| 14 | F | Ideal | 3826 |
| 15 | G | Fair | 314 |
| 16 | G | Good | 871 |
| 17 | G | Very Good | 2299 |
| 18 | G | Premium | 2924 |
| 19 | G | Ideal | 4884 |
| 20 | H | Fair | 303 |
| 21 | H | Good | 702 |
| 22 | H | Very Good | 1824 |
| 23 | H | Premium | 2360 |
| 24 | H | Ideal | 3115 |
| 25 | I | Fair | 175 |
| 26 | I | Good | 522 |
| 27 | I | Very Good | 1204 |
| 28 | I | Premium | 1428 |
| 29 | I | Ideal | 2093 |
| 30 | J | Fair | 119 |
| 31 | J | Good | 307 |
| 32 | J | Very Good | 678 |
| 33 | J | Premium | 808 |
| 34 | J | Ideal | 896 |


Question: How is cut related to color? 2 categorical vars

```
alt.Chart(diamonds_grouped).mark_circle().encode(
  x = 'color:N',
  y = 'cut:N',
  size='N:Q')
```



Question: How is cut related to color? 2 categorical vars

```
alt.Chart(diamonds_grouped).mark_rect().encode(
  x = 'color:N',
  y = 'cut:N',
  color='N:Q')
```



Discussion question: what diamond types are most common?

Two continuous variables

Two continuous variables: roadmap

- `movies` ratings from Rotten Tomatoes and IMDB
- `diamonds`: `carat` vs `price`


movies dataset

```
movies_url = 'https://cdn.jsdelivr.net/npm/vega-datasets@1/data/movies.json'
```

```
movies = pd.read_json(movies_url)
```

Covariation: a first binned scatter plot


```
alt.Chart(movies_url).mark_circle().encode(
  alt.X('Rotten_Tomatoes_Rating:Q', bin=alt.BinParams(maxbins=20)),
  alt.Y('IMDB_Rating:Q', bin=alt.BinParams(maxbins=20)),
)
```

Suffers from overplotting!


use `alt.Size('count()')` to address overplotting

```
xy_size = alt.Chart(movies_url).mark_circle().encode(
    alt.X('Rotten_Tomatoes_Rating:Q', bin=alt.BinParams(maxbins=20)),
    alt.Y('IMDB_Rating:Q', bin=alt.BinParams(maxbins=20)),
    alt.Size('count()')
)
xy_size
```





use `alt.Color('count()')` to address overplotting

```
xy_color = alt.Chart(movies_url).mark_bar().encode(  
    alt.X('Rotten_Tomatoes_Rating:Q', bin=alt.BinParams(maxbins=20)),  
    alt.Y('IMDB_Rating:Q', bin=alt.BinParams(maxbins=20)),  
    alt.Color('count()')  
)  
xy_color
```



Discussion question


`xy_size` | `xy_color`



Compare the *size* and *color*-based 2D histograms above. Which encoding do you think should be preferred? Why?


Question: How is carat related to price? 2 continuous vars

```
alt.Chart(diamonds).mark_point().encode(
  x = 'carat:Q',
  y = 'price:Q'
)
```




Question: How is carat related to price? 2 continuous vars

```
alt.Chart(diamonds).mark_rect().encode(  
  alt.X('carat:Q', bin=alt.Bin(maxbins=70)),  
  alt.Y('price:Q', bin=alt.Bin(maxbins=70)),  
  alt.Color('count()', scale=alt.Scale(scheme='blues')))
```



Question: How is carat related to price? 2 continuous vars

```
alt.Chart(diamonds).mark_boxplot().encode(  
  alt.X('carat:Q', bin=alt.Bin(maxbins=10)),  
  alt.Y('price:Q'))
```




Question: How is carat related to price? 2 continuous vars

```
df = diamonds
df['carat_bin'] = pd.qcut(df['carat'], q=20, labels=(np.arange(1, 21, 1)))

df = df.groupby('carat_bin').agg(
    carat = ('carat', 'mean'),
    price = ('price', 'mean')).reset_index()

alt.Chart(df).mark_point().encode(
    x = 'carat:Q',
    y = 'price:Q'
)
```

/var/folders/9k/556bcdln0hsc_tw0rlx916_00000gn/T/ipykernel_37573/2590590619.py:4: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.



- What it does:
 1. Computes bins using quantiles of x
 2. Computes means of y within each bin
- Called `binscatter` in stata and `binsreg` in R. Doesn't exist yet for Altair, but easy to code up yourself

Discussion question – “How is carat related to price?”

Review the `mark_rect()`, `mark_boxplot()`, and `binscatter` plots

- headline? (aka the main message)
- sub-messages? (other information one can learn beyond the main message)

Summary: Exploring covariation

| Scenario | Functions |
|-------------------------------------|----------------------------------|
| Categorical and continuous variable | <code>mark_boxplot()</code> |
| | <code>transform_density()</code> |
| Two categorical variables | <code>size</code> |
| | <code>color</code> |
| Two continuous variables | <code>alt.Size('count()')</code> |

| Scenario | Functions |
|----------|-----------------------------------|
| | <code>alt.Color('count()')</code> |
| | <code>mark_boxplot()</code> |
| | <code>binscatter</code> |

Do-pair-share


We are now going to transition from making plots to teach **ourselves** to making plots for an audience.

Are penguins getting heavier (`body_mass_g`) over time?

Bonus: what is the headline of your plot and what are the sub-messages?

Do-pair-share solution I


```
alt.Chart(penguins).mark_bar().encode(
  alt.Y('average(body_mass_g):Q', scale=alt.Scale(zero=False)),
  alt.X('year:N'),
  alt.Color('year:N')
)
```



This does answers the question, albeit in the most simple/boring way possible.

Do-pair-share solution II

```
alt.Chart(penguins).transform_density(
    'body_mass_g',
    groupby=['year'],
    as_ = ['body_mass_g', 'density']
).mark_line().encode(
    x = 'body_mass_g:Q',
    y = 'density:Q',
    color='year:N'
)
```



- Headline: 2007 is lightest, 2008 is heaviest
- Sub-messages
 1. Similar shares of penguins above 5,000 grams in 2008 and 2009
 2. Average weight is higher in 2008 because 2009 has more lightweight penguins

Meta comment: iterating on plot design

“Make dozens of plots” – Quoctrung Bui, former 30535 guest lecturer and former Harris data viz instructor

What does he mean?

- The first plot you make will never be the one you should show
- As a rule of thumb, you should try out at least three different plotting concepts (marks)
- Within each concept, you will need to try out several different encodings

